

Project Proposal — Forecasting One Piece TCG Swiss Matches

Kiana Kiser · MLOps HS26 · I.BA_MLOPS · Repository (public): github.com/kianakiser/optcg-match-forecast

1 Problem statement

What and for whom. For every Swiss match at an online One Piece TCG tournament, predict the probability the seat-1 player wins, from both registered decklists and both pilots' results at strictly earlier events. For competitive players and deck-builders.

Horizon. One match, priced at pairing time and resolved within the hour it is played. The model is given only what is known before the first card is played — the two registered 50-card lists and both pilots' results at strictly earlier events — and predicts that match's outcome. It is served on demand through a web UI (a player picks two decks and gets a probability), and the same model is scored nightly against every match that resolved since the previous run. The provider lists a tournament only once it has finished, so there is no in-progress feed to predict against in real time; the honest framing is on-demand serving plus nightly scoring, not live in-tournament inference. The reported winner is the label and is not computable from any feature.

Scope. Online Swiss rounds, predominantly best-of-one, at 32+ entrant events on play.limitlesstcg.com. Top cut, byes, ties and unresolved seats excluded — including two events whose single-elimination brackets are mislabelled as Swiss. One corpus is used throughout: the 198 events ingested to 2026-09-18, **28,942 decided Swiss matches** over 12,166 entrants. 82.8% of events and 80.2% of matches come from one recurring series, so claims cover online play only.

Success criterion. Baseline: a coin flip — Brier 0.2500, log loss 0.6931 — because the label is balanced by construction: seat 1 wins 50.56% of those 28,942 matches (Wilson 95% CI 49.98–51.14%), leaving no majority-class rule and no rare positive class. Success = pooled over ≥ 6 held-out 28-day windows **split by event date, never by row** ($\geq 8,000$ matches): Brier ≤ 0.2490 with the skill interval (0.2500 – Brier, whole-event bootstrap) excluding zero, and calibration gaps ≤ 3 points in every 5-point favourite bucket with $n \geq 250$. Reproducing the existing archetype model out-of-time gives 0.2465; 0.2490 is the bar every six-window stretch in 14 months clears.

2 Originality & motivation

[Kiana writes this — see [notes/section2_worksheet.md](#). 150 words, three things: why you (you already run this pipeline); why the problem (the published win rate conflates deck with pilot, and you can show the gap); and the checked-lists sentence. Both lists ARE checked: [mlops-lab.ch](#) FS26 has no card game, but KTH 2026 does have chess/NHL/football match predictors — so differentiate on structure, not on novelty.]

3 Data source & features

Scope. Training uses the full international online field. Swiss players are 4.0% of entrants (166 distinct players, 1,965 matches involving at least one), reported as a secondary evaluation slice rather than a training restriction: 525 Swiss-vs-Swiss matches would be far too few to fit.

Matches come from Limitless (play.limitlesstcg.com) via its public JSON API: `/api/tournaments,{id}/standings,{id}/pairings` — no auth, no scraping. GitHub Actions ingests daily at 06:07 UTC; finished events are immutable and fetched once. At 2026-09-18 the corpus holds **28,942 decided Swiss matches** over 198 events and 12,166 entrants, **11,575 of them (95.1%) with a full 50-card list**, growing ≈ 570 matches and 3.5 events weekly. The source is live, not an archive: seven in-scope events existed in the index but not in the corpus when this was written, the newest a day old.

Label. `pairings[].winner`, the username the platform records from the reported result; no decklist rule yields it. Seat 1 wins 50.56% on that same corpus, so no rare positive class exists. Scarcity is in the archetype tail: a new set every 2–3 months introduces leaders with no history at all, and 3–5 of the top 10 archetypes turn over at each release. Those rows stay in: the service returns 0.5 with an explicit coverage flag rather than guessing, because a leader-attribute backoff measured worse than a coin flip.

Features, per side and differenced: `leader_id`, `leader_life`, `leader_power`, `leader_color_count`; deck aggregates `counter_2k_copies`, `avg_cost`, `high_curve_copies`, `big_body_copies`, `event_card_copies`, `distinct_cards`; and `player_prior_winrate` and `archetype_prior_winrate` over strictly earlier events.

Leakage. `placing`, `record` and `drop` are outcomes of the predicted event and are dropped at ingest: `placing` is null on 43.0% of 12,166 raw entries, all of which carry `drop`, so filtering `placing IS NOT NULL` would condition on finishing. Splits are rolling-origin by event date, never by row (one event's matches share decks and pilots): train up to T, test (T, T+28d]. The ingest must also persist the event date, which only the tournament index holds.

4 System design




Figure 1: The three FTI pipelines. They are decoupled — none calls another; they meet only at the feature store and the model registry.

Core (ships first). A GitHub Actions job runs daily at 06:07 UTC, matching the source's ≈ 3.5 events a week. It pulls the Limitless API, drops placing, record and drop physically at the ingest boundary, persists each event's date from the tournament index — the pairings payload carries none, so without it no out-of-time split is possible — lands raw JSON in GCS and writes features to Hopsworks. Training reads the feature view, splits by event date, evaluates out of time, and logs every run to Weights & Biases, and registers the winner in the Hopsworks model registry under the alias *champion*. Cloud Run loads whatever currently carries that alias and serves win probabilities to the UI, logging each prediction so it can be scored once the match resolves; serving never calls the provider.

Stack. *Hopsworks*: point-in-time joins, so player history cannot include the event being predicted, and one hosted service covers both the feature store and the registry — promotion moves an alias rather than redeploying. *Weights & Biases*: hosted run tracking, so there is no MLflow server to operate for a project this size. *GitHub Actions*: schedules beside code and CI, and sole holder of the provider client — throttled to the published 50 requests per 5 minutes, finished events cached write-once. *Cloud Run*: container deploy, scales to zero between events. *GCS*: immutable landing zone, so features rebuild without re-fetching.

Optional stretch, outside the core: decklist drift monitoring, a second feed to dilute the $\approx 82\%$ -of-matches single-organiser concentration, and a calibration dashboard.

The repository is public: github.com/kianakiser/optcg-match-forecast.